

# Package ‘ragnar’

August 19, 2025

**Title** Retrieval-Augmented Generation (RAG) Workflows

**Version** 0.2.1

**Description** Provides tools for implementing Retrieval-Augmented Generation (RAG) workflows with Large Language Models (LLM). Includes functions for document processing, text chunking, embedding generation, storage management, and content retrieval. Supports various document types and embedding providers ('Ollama', 'OpenAI'), with 'DuckDB' as the default storage backend. Integrates with the 'ellmer' package to equip chat objects with retrieval capabilities. Designed to offer both sensible defaults and customization options with transparent access to intermediate outputs. For a review of retrieval-augmented generation methods, see Gao et al. (2023) ``Retrieval-Augmented Generation for Large Language Models: A Survey" <doi:10.48550/arXiv.2312.10997>.

**License** MIT + file LICENSE

**URL** <https://ragnar.tidyverse.org/>, <https://github.com/tidyverse/ragnar>

**BugReports** <https://github.com/tidyverse/ragnar/issues>

**Depends** R (>= 4.3.0)

**Imports** blob, cli, commonmark, curl, DBI, dbplyr, dplyr, duckdb (>= 1.3.1), glue, httr2, methods, reticulate (>= 1.42.0), rlang (>= 1.1.0), rvest, S7, stringi, tidyr, vctrs, withr, xml2

**Suggests** connectcreds, ellmer (>= 0.3.0), gargle, knitr, lifecycle, pandoc, paws.common, rmarkdown, shiny, stringr, tibble, testthat (>= 3.0.0)

**VignetteBuilder** knitr

**Config/Needs/website** tidyverse/tidytemplate, rmarkdown

**Config/testthat/edition** 3

**Encoding** UTF-8

**RoxygenNote** 7.3.2

**NeedsCompilation** yes

**Author** Tomasz Kalinowski [aut, cre],  
Daniel Falbel [aut],  
Posit Software, PBC [cph, fnd] (ROR: <<https://ror.org/03wc8by49>>)  
**Maintainer** Tomasz Kalinowski <tomasz@posit.co>  
**Repository** CRAN  
**Date/Publication** 2025-08-19 21:40:02 UTC

Contents

chunks_deoverlap . . . . .	2
embed_bedrock . . . . .	3
embed_databricks . . . . .	4
embed_google_gemini . . . . .	4
embed_ollama . . . . .	6
MarkdownDocument . . . . .	8
MarkdownDocumentChunks . . . . .	9
markdown_chunk . . . . .	10
ragnar_chunks_view . . . . .	12
ragnar_find_links . . . . .	12
ragnar_register_tool_retrieve . . . . .	13
ragnar_retrieve . . . . .	14
ragnar_retrieve_bm25 . . . . .	15
ragnar_retrieve_vss . . . . .	16
ragnar_store_build_index . . . . .	18
ragnar_store_create . . . . .	18
ragnar_store_insert . . . . .	21
ragnar_store_inspect . . . . .	22
read_as_markdown . . . . .	22
<b>Index</b>	<b>26</b>

---

chunks_deoverlap	<i>Merge overlapping chunks in retrieved results</i>
------------------	--

---

Description

Groups and merges overlapping text chunks from the same origin in the retrieval results.

Usage

chunks\_deoverlap(store, chunks)

Arguments

- store            A RagnarStore object. Must have @version == 2.
- chunks          A [tibble](#) of retrieved chunks, such as the output of [ragnar\\_retrieve\(\)](#).

Details

When multiple retrieved chunks from the same origin have overlapping character ranges, this function combines them into a single non-overlapping region.

Value

A [tibble](#) of de-overlapped chunks.

---

embed_bedrock	<i>Embed text using a Bedrock model</i>
---------------	---

---

Description

Embed text using a Bedrock model

Usage

```
embed_bedrock(x, model, profile, api_args = list())
```

Arguments

x	x can be: <ul style="list-style-type: none"><li>• A character vector, in which case a matrix of embeddings is returned.</li><li>• A data frame with a column named <code>text</code>, in which case the dataframe is returned with an additional column named <code>embedding</code>.</li><li>• Missing or NULL, in which case a function is returned that can be called to get embeddings. This is a convenient way to partial in additional arguments like <code>model</code>, and is the most convenient way to produce a function that can be passed to the <code>embed</code> argument of <code>ragnar_store_create()</code>.</li></ul>
model	Currently only Cohere.ai and Amazon Titan models are supported. There are no guardrails for the kind of model that is used, but the model must be available in the AWS region specified by the profile. You may look for available models in the Bedrock Model Catalog
profile	AWS profile to use.
api_args	Additional arguments to pass to the Bedrock API. Depending on the model, you might be able to provide different parameters. Check the documentation for the model you are using in the <a href="#">Bedrock user guide</a> .

Value

If `x` is missing returns a function that can be called to get embeddings. If `x` is not missing, a matrix of embeddings with 1 row per input string, or a dataframe with an `'embedding'` column.

See Also

[embed\\_ollama\(\)](#)

---

embed_databricks	<i>Embed text using a Databricks model</i>
------------------	--

---

**Description**

`embed_databricks()` gets embeddings for text using a model hosted in a Databricks workspace. It relies on the **ellmer** package for managing Databricks credentials. See `ellmer::chat_databricks` for more on supported modes of authentication.

**Usage**

```
embed_databricks(  
  x,  
  workspace = databricks_workspace(),  
  model = "databricks-bge-large-en",  
  batch_size = 512L  
)
```

**Arguments**

x	<p>x can be:</p> <ul style="list-style-type: none"><li>• A character vector, in which case a matrix of embeddings is returned.</li><li>• A data frame with a column named <code>text</code>, in which case the dataframe is returned with an additional column named <code>embedding</code>.</li><li>• Missing or <code>NULL</code>, in which case a function is returned that can be called to get embeddings. This is a convenient way to partial in additional arguments like <code>model</code>, and is the most convenient way to produce a function that can be passed to the <code>embed</code> argument of <code>ragnar_store_create()</code>.</li></ul>
workspace	The URL of a Databricks workspace, e.g. <code>"https://example.cloud.databricks.com"</code> . Will use the value of the environment variable <code>DATABRICKS_HOST</code> , if set.
model	The name of a text embedding model.
batch_size	split x into batches when embedding. Integer, limit of strings to include in a single request.

---

embed_google_gemini	<i>Embed using Google Vertex API platform</i>
---------------------	---

---

**Description**

Embed using Google Vertex API platform

**Usage**

```

embed_google_gemini(
  x,
  model = "gemini-embedding-001",
  base_url = "https://generativelanguage.googleapis.com/v1beta",
  api_key = get_envvar("GEMINI_API_KEY"),
  dims = NULL,
  task_type = "RETRIEVAL_QUERY",
  batch_size = 20L
)

embed_google_vertex(
  x,
  model,
  location,
  project_id,
  task_type = "RETRIEVAL_QUERY"
)

```

**Arguments**

x	<p>x can be:</p> <ul style="list-style-type: none"> <li>• A character vector, in which case a matrix of embeddings is returned.</li> <li>• A data frame with a column named <code>text</code>, in which case the dataframe is returned with an additional column named <code>embedding</code>.</li> <li>• Missing or <code>NULL</code>, in which case a function is returned that can be called to get embeddings. This is a convenient way to partial in additional arguments like <code>model</code>, and is the most convenient way to produce a function that can be passed to the <code>embed</code> argument of <code>ragnar_store_create()</code>.</li> </ul>
model	Character specifying the embedding model. See supported models in <a href="#">Text embeddings API</a>
base_url	string, url where the service is available.
api_key	resolved using env var <code>GEMINI_API_KEY</code>
dims	An integer, can be used to truncate the embedding to a specific size.
task_type	<p>Used to convey intended downstream application to help the model produce better embeddings. If left blank, the default used is <code>"RETRIEVAL_QUERY"</code>.</p> <ul style="list-style-type: none"> <li>• <code>"RETRIEVAL_QUERY"</code></li> <li>• <code>"RETRIEVAL_DOCUMENT"</code></li> <li>• <code>"SEMANTIC_SIMILARITY"</code></li> <li>• <code>"CLASSIFICATION"</code></li> <li>• <code>"CLUSTERING"</code></li> <li>• <code>"QUESTION_ANSWERING"</code></li> <li>• <code>"FACT_VERIFICATION"</code></li> <li>• <code>"CODE_RETRIEVAL_QUERY"</code> For more information about task types, see <a href="#">Choose an embeddings task type</a>.</li> </ul>

batch_size	split x into batches when embedding. Integer, limit of strings to include in a single request.
location	Location, e.g. us-east1, me-central1, africa-south1.
project_id	Project ID.

Functions

- embed\_google\_gemini(): Use the Gemini API to create embeddings.

Examples

```
embed_google_gemini("hello world")

## Not run:
embed_google_vertex(
  "hello world",
  model="gemini-embedding-001",
  project = "<your-project-id>",
  location = "us-central1"
)

## End(Not run)
```

---

embed_ollama	<i>Embed Text</i>
--------------	-------------------

---

Description

Embed Text

Usage

```
embed_ollama(
  x,
  base_url = "http://localhost:11434",
  model = "snowflake-arctic-embed2:568m",
  batch_size = 10L
)

embed_openai(
  x,
  model = "text-embedding-3-small",
  base_url = "https://api.openai.com/v1",
  api_key = get_envvar("OPENAI_API_KEY"),
  dims = NULL,
  user = get_user(),
  batch_size = 20L
)
```

```

embed_lm_studio(
  x,
  model,
  base_url = "http://localhost:1234/v1",
  api_key = "lm-studio",
  dims = NULL,
  user = get_user(),
  batch_size = 20L
)

```

## Arguments

x	<p>x can be:</p> <ul style="list-style-type: none"> <li>• A character vector, in which case a matrix of embeddings is returned.</li> <li>• A data frame with a column named <code>text</code>, in which case the dataframe is returned with an additional column named <code>embedding</code>.</li> <li>• Missing or <code>NULL</code>, in which case a function is returned that can be called to get embeddings. This is a convenient way to partial in additional arguments like <code>model</code>, and is the most convenient way to produce a function that can be passed to the <code>embed</code> argument of <code>ragnar_store_create()</code>.</li> </ul>
base_url	string, url where the service is available.
model	string; model name
batch_size	split x into batches when embedding. Integer, limit of strings to include in a single request.
api_key	resolved using env var <code>OPENAI_API_KEY</code>
dims	An integer, can be used to truncate the embedding to a specific size.
user	User name passed via the API.

## Value

If `x` is a character vector, then a numeric matrix is returned, where `nrow = length(x)` and `ncol = <model-embedding-size>`.  
 If `x` is a data.frame, then a new embedding matrix "column" is added, containing the matrix described in the previous sentence.

A matrix of embeddings with 1 row per input string, or a dataframe with an 'embedding' column.

## Functions

- `embed_lm_studio()`: Embed Text using LMStudio. Identical to `embed_openai()` but with suitable defaults for LMStudio.

## Examples

```

text <- c("a chunk of text", "another chunk of text", "one more chunk of text")
## Not run:
text |>
  embed_ollama() |>

```

```

    str()

text |>
  embed_openai() |>
    str()

## End(Not run)

```

---

MarkdownDocument	<i>Markdown documents</i>
------------------	---------------------------

---

## Description

MarkdownDocument represents a complete Markdown document stored as a single character string. The constructor normalizes text by collapsing lines and ensuring UTF-8 encoding, so downstream code can rely on a consistent format.

[read\\_as\\_markdown\(\)](#) is the recommended way to create a MarkdownDocument. The constructor itself is exported only so advanced users can construct one by other means when needed.

## Arguments

text	[string] Markdown text.
origin	[string] Optional source path or URL. Defaults to the "origin" attribute of text, if present, otherwise NULL.

## Value

An S7 object that inherits from MarkdownDocument, which is a length 1 string of markdown text with an @origin property.

## Examples

```

md <- MarkdownDocument(
  "# Title\n\nSome text.",
  origin = "example.md"
)
md

```



---

 MarkdownDocumentChunks

*Markdown documents chunks*


---

## Description

MarkdownDocumentChunks stores information about candidate chunks in a Markdown document. It is a tibble with three required columns:

- `start`, `end` — integers. These are character positions (1-based, inclusive) in the source MarkdownDocument, so that `substring(md, start, end)` yields the chunk text. Ranges can overlap.
- `context` — character. A general-purpose field for adding context to a chunk. This column is combined with text to augment chunk content when generating embeddings with `ragnar_store_insert()`, and is also returned by `ragnar_retrieve()`. Keep in mind that when chunks are deoverlapped (in `ragnar_retrieve()` or `chunks_deoverlap()`), only the context value from the first chunk is kept. `markdown_chunk()` by default populates this column with all the markdown headings that are in-scope at the chunk start position.

Additional columns can be included.

The original document is available via the `@document` property.

For normal use, chunk a Markdown document with `markdown_chunk()`; the class constructor itself is exported only so advanced users can generate or tweak chunks by other means.

## Arguments

<code>chunks</code>	A data frame containing <code>start</code> , <code>end</code> , and <code>context</code> columns, and optionally other columns.
<code>document</code>	A MarkdownDocument.

## Value

An S7 object that inherits from MarkdownDocumentChunks, which is also a tibble.

## See Also

[MarkdownDocument\(\)](#)

## Examples

```
doc_text <- "## A\n\nB\n\n## C\n\nD" # can be readLines() output, etc.
doc <- MarkdownDocument(doc_text, origin = "some/where")
chunk_positions <- tibble::tibble(
  start = c(1L, 9L),
  end = c(8L, 15L),
  context = c("", "## A"),
  text = substring(doc, start, end)
```

```
)
chunks <- MarkdownDocumentChunks(chunk_positions, doc)
identical(chunks@document, doc)
```

---

markdown\_chunk

*Chunk a Markdown document*


---

## Description

markdown\_chunk() splits a single Markdown string into shorter optionally overlapping chunks while nudging cut points to the nearest sensible boundary (heading, paragraph, sentence, line, word, or character). It returns a tibble recording the character ranges, headings context, and text for each chunk.

## Usage

```
markdown_chunk(
  md,
  target_size = 1600L,
  target_overlap = 0.5,
  ...,
  max_snap_dist = target_size * (1 - target_overlap)/3,
  segment_by_heading_levels = integer(),
  context = TRUE,
  text = TRUE
)
```

## Arguments

md	A MarkdownDocument, or a length-one character vector containing Markdown.
target_size	Integer. Target chunk size in characters. Default: 1600 ( $\approx$ 400 tokens, or 1 page of text). Actual chunk size may differ from the target by up to $2 * \text{max\_snap\_dist}$ . When set to NULL, NA or Inf and used with segment_by_heading_levels, chunk size is unbounded and each chunk corresponds to a segment.
target_overlap	Numeric in $[0, 1)$ . Fraction of desired overlap between successive chunks. Default: 0.5. Even when 0, some overlap can occur because the last chunk is anchored to the document end.
...	These dots are for future extensions and must be empty.
max_snap_dist	Integer. Furthest distance (in characters) a cut point may move to reach a semantic boundary. Defaults to one third of the stride size between target chunk starts. Chunks that end up on identical boundaries are merged.
segment_by_heading_levels	Integer vector with possible values 1:6. Headings at these levels are treated as segment boundaries; chunking is performed independently for each segment. No chunk will overlap a segment boundary, and any future deoverlapping will not combine segments. Each segment will have a chunk that starts at the segment start and a chunk that ends at the segment end (these may be the same chunk or overlap substantially if the segment is short). Default: disabled.

context	Logical. Add a context column containing the Markdown headings in scope at each chunk start. Default: TRUE.
text	Logical. If TRUE, include a text column with the chunk contents. Default: TRUE.

### Value

A [MarkdownDocumentChunks](#) object, which is a tibble (data.frame) with with columns `start` `end`, and optionally `context` and `text`. It also has a `@document` property, which is the input `md` document (potentially normalized and converted to a [MarkdownDocument](#)).

### See Also

[ragnar\\_chunks\\_view\(\)](#) to interactively inspect the output of `markdown_chunk()`. See also [MarkdownDocumentChunks\(\)](#) and [MarkdownDocument\(\)](#), where the input and return value of `markdown_chunk()` are described more fully.

### Examples

```
md <- "
# Title

## Section 1

Some text that is long enough to be chunked.

A second paragraph to make the text even longer.

## Section 2

More text here.

### Section 2.1

Some text under a level three heading.

#### Section 2.1.1

Some text under a level four heading.

## Section 3

Even more text here.
"

markdown_chunk(md, target_size = 40)
markdown_chunk(md, target_size = 40, target_overlap = 0)
markdown_chunk(md, target_size = NA, segment_by_heading_levels = c(1, 2))
markdown_chunk(md, target_size = 40, max_snap_dist = 100)
```

---

ragnar_chunks_view	<i>View chunks with the store inspector</i>
--------------------	---

---

### Description

Visualize chunks read by `ragnar_read()` for quick inspection. Helpful for inspecting the results of chunking and reading while iterating on the ingestion pipeline.

### Usage

```
ragnar_chunks_view(chunks)
```

### Arguments

chunks	A data frame containing a few chunks.
--------	---------------------------------------

---

ragnar_find_links	<i>Find links on a page</i>
-------------------	-----------------------------

---

### Description

Find links on a page

### Usage

```
ragnar_find_links(
  x,
  depth = 0L,
  children_only = TRUE,
  progress = TRUE,
  ...,
  url_filter = identity,
  validate = FALSE
)
```

### Arguments

x	URL, HTML file path, or XML document. For Markdown, convert to HTML using <code>commonmark::markdown_html()</code> first.
depth	Integer specifying how many levels deep to crawl for links. When <code>depth &gt; 0</code> , the function will follow child links (links with x as a prefix) and collect links from those pages as well.
children_only	Logical or string. If TRUE, returns only child links (those having x as a prefix). If FALSE, returns all links found on the page. Note that regardless of this setting, only child links are followed when <code>depth &gt; 0</code> .

progress	Logical, draw a progress bar if depth > 0.
...	Currently unused. Must be empty.
url_filter	A function that takes a character vector of URL's and may subset them to return a smaller list. This can be useful for filtering out URL's by rules different them children_only which only checks the prefix.
validate	Default is FALSE. If TRUE sends a HEAD request for each link and removes those that are not accessible. Requests are sent in parallel using <a href="#">httr2::req_perform_parallel()</a> .

**Value**

A character vector of links on the page.

**Examples**

```
## Not run:
ragnar_find_links("https://r4ds.hadley.nz/base-R.html")
ragnar_find_links("https://ellmer.tidyverse.org/")
ragnar_find_links("https://ellmer.tidyverse.org/", depth = 2)
ragnar_find_links("https://ellmer.tidyverse.org/", depth = 2, children_only = FALSE)
ragnar_find_links(
  paste0("https://github.com/Snowflake-Labs/sfquickstarts/",
        "tree/master/site/sfguides/src/build_a_custom_model_for_anomaly_detection"),
  children_only = "https://github.com/Snowflake-Labs/sfquickstarts",
  depth = 1
)

## End(Not run)
```

---

```
ragnar_register_tool_retrieve
```

*Register a 'retrieve' tool with ellmer*

---

**Description**

Register a 'retrieve' tool with ellmer

**Usage**

```
ragnar_register_tool_retrieve(
  chat,
  store,
  store_description = "the knowledge store",
  ...,
  name = NULL,
  title = NULL
)
```

**Arguments**

chat	a <code>ellmer:::Chat</code> object.
store	a string of a store location, or a <code>RagnarStore</code> object.
store_description	Optional string, used for composing the tool description.
...	arguments passed on to <code>ragnar_retrieve()</code> .
name, title	Optional tool function name and title. By default, <code>store@name</code> and <code>store@title</code> will be used if present. The tool name must be a valid R function name and should be unique with the tools registered with the <code>ellmer::Chat</code> object. <code>title</code> is used for user-friendly display.

**Value**

chat, invisibly.

**Examples**

```
system_prompt <- stringr::str_squish("
  You are an expert assistant in R programming.
  When responding, you first quote relevant material from books or documentation,
  provide links to the sources, and then add your own context and interpretation.
")
chat <- ellmer::chat_openai(system_prompt, model = "gpt-4.1")

store <- ragnar_store_connect("r4ds.ragnar.duckdb")
ragnar_register_tool_retrieve(chat, store)
chat$chat("How can I subset a dataframe?")
```

---

ragnar_retrieve	<i>Retrieve chunks from a RagnarStore</i>
-----------------	---

---

**Description**

Combines both vss and bm25 search and returns the union of chunks retrieved by both methods.

**Usage**

```
ragnar_retrieve(store, text, top_k = 3L, ..., deoverlap = TRUE)
```

**Arguments**

store	A <code>RagnarStore</code> object returned by <code>ragnar_store_connect()</code> or <code>ragnar_store_create()</code> .
text	Character. Query string to match.
top_k	Integer. Number of nearest entries to find per method.
...	Additional arguments passed to the lower-level retrieval functions.
deoverlap	Logical. If TRUE (default) and <code>store@version == 2</code> , overlapping chunks are merged with <code>chunks_deoverlap()</code> .

**Value**

A tibble of retrieved chunks. Each row represents a chunk and always contains a text column.

**Note**

The results are not re-ranked after identifying the unique values.

**See Also**

Other ragnar\_retrieve: [ragnar\\_retrieve\\_bm25\(\)](#), [ragnar\\_retrieve\\_vss\(\)](#), [ragnar\\_retrieve\\_vss\\_and\\_bm25\(\)](#)

**Examples**

```
## Build a small store with categories
store <- ragnar_store_create(
  embed = \(x) ragnar::embed_openai(x, model = "text-embedding-3-small"),
  extra_cols = data.frame(category = character()),
  version = 1 # store text chunks directly
)

ragnar_store_insert(
  store,
  data.frame(
    category = c(rep("pets", 3), rep("dessert", 3)),
    text      = c("playful puppy", "sleepy kitten", "curious hamster",
                  "chocolate cake", "strawberry tart", "vanilla ice cream")
  )
)
ragnar_store_build_index(store)

# Top 3 chunks without filtering
ragnar_retrieve(store, "sweet")

# Combine filter with similarity search
ragnar_retrieve(store, "sweet", filter = category == "dessert")
```

---

`ragnar_retrieve_bm25`    *Retrieves chunks using the BM25 score*

---

**Description**

BM25 refers to Okapi Best Matching 25. See [doi:10.1561/15000000019](https://doi.org/10.1561/15000000019) for more information.

**Usage**

```
ragnar_retrieve_bm25(
  store,
  text,
  top_k = 3L,
  ...,
  k = 1.2,
  b = 0.75,
  conjunctive = FALSE,
  filter
)
```

**Arguments**

store	A RagnarStore object returned by <code>ragnar_store_connect()</code> or <code>ragnar_store_create()</code> .
text	String, the text to search for.
top_k	Integer. Number of nearest entries to find per method.
...	Additional arguments passed to the lower-level retrieval functions.
k, b	$k_1$ and $b$ parameters in the Okapi BM25 retrieval method.
conjunctive	Whether to make the query conjunctive i.e., all terms in the query string must be present in order for a chunk to be retrieved.
filter	Optional. A filter expression evaluated with <code>dplyr::filter()</code> .

**See Also**

Other ragnar\_retrieve: [ragnar\\_retrieve\(\)](#), [ragnar\\_retrieve\\_vss\(\)](#), [ragnar\\_retrieve\\_vss\\_and\\_bm25\(\)](#)

---

ragnar_retrieve_vss	<i>Vector Similarity Search Retrieval</i>
---------------------	---

---

**Description**

Computes a similarity measure between the query and the document embeddings and uses this similarity to rank and retrieve document chunks.

**Usage**

```
ragnar_retrieve_vss(
  store,
  query,
  top_k = 3L,
  ...,
  method = "cosine_distance",
  query_vector = store@embed(query),
  filter
)
```



**Arguments**

store	A RagnarStore object returned by <code>ragnar_store_connect()</code> or <code>ragnar_store_create()</code> .
query	Character. The query string to embed and use for similarity search.
top_k	Integer. Maximum number of document chunks to retrieve. Defaults to 3.
...	Additional arguments passed to methods.
method	Character. Similarity method to use: "cosine_distance", "euclidean_distance", or "negative_inner_product". Defaults to "cosine_distance".
query_vector	Numeric vector. The embedding for query. Defaults to <code>store@embed(query)</code> .
filter	Optional. A filter expression evaluated with <code>dplyr::filter()</code> .

**Details**

Supported methods:

- **cosine\_distance** – cosine of the angle between two vectors.
- **euclidean\_distance** – L2 distance between vectors.
- **negative\_inner\_product** – negative sum of element-wise products.

If `filter` is supplied, the function first performs the similarity search, then applies the filter in an outer SQL query. It uses the HNSW index when possible and falls back to a sequential scan for large result sets or filtered queries.

**Value**

A tibble with the `top_k` retrieved chunks, ordered by `metric_value`.

**Note**

The results are not re-ranked after identifying the unique values.

**See Also**

Other `ragnar_retrieve`: [ragnar\\_retrieve\(\)](#), [ragnar\\_retrieve\\_bm25\(\)](#), [ragnar\\_retrieve\\_vss\\_and\\_bm25\(\)](#)

**Examples**

```
## Build a small store with categories
store <- ragnar_store_create(
  embed = \(x) ragnar::embed_openai(x, model = "text-embedding-3-small"),
  extra_cols = data.frame(category = character()),
  version = 1 # store text chunks directly
)

ragnar_store_insert(
  store,
  data.frame(
    category = c(rep("pets", 3), rep("dessert", 3)),
    text      = c("playful puppy", "sleepy kitten", "curious hamster",
```

```

        "chocolate cake", "strawberry tart", "vanilla ice cream")
    )
)
ragnar_store_build_index(store)

# Top 3 chunks without filtering
ragnar_retrieve(store, "sweet")

# Combine filter with similarity search
ragnar_retrieve(store, "sweet", filter = category == "dessert")

```

---

ragnar\_store\_build\_index

*Build a Ragnar Store index*


---

### Description

A search index must be built before calling `ragnar_retrieve()`. If additional entries are added to the store with `ragnar_store_insert()`, `ragnar_store_build_index()` must be called again to rebuild the index.

### Usage

```
ragnar_store_build_index(store, type = c("vss", "fts"))
```

### Arguments

store	a RagnarStore object
type	The retrieval search type to build an index for.

### Value

store, invisibly.

---

ragnar\_store\_create

*Create and connect to a vector store*


---

### Description

Create and connect to a vector store

**Usage**

```
ragnar_store_create(
  location = ":memory:",
  embed = embed_ollama(),
  ...,
  embedding_size = ncol(embed("foo")),
  overwrite = FALSE,
  extra_cols = NULL,
  name = NULL,
  title = NULL,
  version = 2
)

ragnar_store_connect(location, ..., read_only = TRUE)
```

**Arguments**

location	filepath, or :memory:. Location can also be a database name specified with md:dbname, in this case the database will be created in MotherDuck after a connection is established.
embed	A function that is called with a character vector and returns a matrix of embeddings. Note this function will be serialized and then deserialized in new R sessions, so it cannot reference to any objects in the global or parent environments. Make sure to namespace all function calls with :. If additional R objects must be available in the function, you can optionally supply a carrier::crate() with packaged data. It can also be NULL for stores that don't need to embed their texts, for example, if only using FTS algorithms such as <a href="#">ragnar_retrieve_bm25()</a> .
...	unused; must be empty.
embedding_size	integer
overwrite	logical, what to do if location already exists
extra_cols	A zero row data frame used to specify additional columns that should be added to the store. Such columns can be used for adding additional context when retrieving. See the examples for more information. <a href="#">vctrs::vec_cast()</a> is used to consistently perform type checks and casts when inserting with <a href="#">ragnar_store_insert()</a> .
name	A unique name for the store. Must match the <code>^[a-zA-Z0-9_-]+\$</code> regex. Used by <a href="#">ragnar_register_tool_retrieve()</a> for registering tools.
title	A title for the store, used by <a href="#">ragnar_register_tool_retrieve()</a> when the store is registered with an <a href="#">ellmer::Chat</a> object.
version	integer. The version of the store to create. See details.
read_only	logical, whether the returned connection can be used to modify the store.

**Details****Store versions:**

**Version 2 – documents with chunk ranges** (default)

With `version = 2`, ragnar stores each document once and records the start and end positions of its chunks. This provides strong support for overlapping chunk ranges with de-overlapping at retrieval, and generally allows retrieving arbitrary ranges from source documents, but does not support modifying chunks directly before insertion. Chunks can be augmented via the `context` field and with additional fields passed to `extra_cols`. The easiest way to prepare chunks for `version = 2` is with `read_as_markdown()` and `markdown_chunk()`.

### Version 1 – flat chunks

With `version = 1`, ragnar keeps all chunks in a single table. This lets you easily modify chunk text before insertion. However, dynamic rechunking (de-overlapping) or extracting arbitrary ranges from source documents is not supported, since the original full documents are no longer available. Chunks can be augmented by modifying the chunk text directly (e.g., with `glue()`). Additionally, if you intend to call `ragnar_store_update()`, it is your responsibility to provide `rlang::hash(original_full_document)` with each chunk. The easiest way to prepare chunks for `version = 1` is with `ragnar_read()` and `ragnar_chunk()`.

## Value

a `RagnarStore` object

## Examples

```
# A store with a dummy embedding
store <- ragnar_store_create(
  embed = \(x) matrix(stats::runif(10), nrow = length(x), ncol = 10),
  version = 1
)
ragnar_store_insert(store, data.frame(text = "hello"))

# A store with a schema. When inserting into this store, users need to
# provide an `area` column.
store <- ragnar_store_create(
  embed = \(x) matrix(stats::runif(10), nrow = length(x), ncol = 10),
  extra_cols = data.frame(area = character()),
  version = 1
)
ragnar_store_insert(store, data.frame(text = "hello", area = "rag"))

# If you already have a data.frame with chunks that will be inserted into
# the store, you can quickly create a suitable store with `vec_ptype()`:
chunks <- data.frame(text = letters, area = "rag")
store <- ragnar_store_create(
  embed = \(x) matrix(stats::runif(10), nrow = length(x), ncol = 10),
  extra_cols = vctrs::vec_ptype(chunks),
  version = 1
)
ragnar_store_insert(store, chunks)

# version = 2 (the default) has support for deoverlapping
store <- ragnar_store_create(
  # if embed = NULL, then only bm25 search is used (not vss)
  embed = NULL
)
```

```

)
doc <- MarkdownDocument(
  paste0(letters, collapse = ""),
  origin = "/some/where"
)
chunks <- markdown_chunk(doc, target_size = 3, target_overlap = 2 / 3)
chunks$context <- substring(chunks$text, 1, 1)
chunks
ragnar_store_insert(store, chunks)
ragnar_store_build_index(store)

ragnar_retrieve(store, "abc bcd xyz", deoverlap = FALSE)
ragnar_retrieve(store, "abc bcd xyz", deoverlap = TRUE)

```

---

`ragnar_store_insert`     *Inserts or updates chunks in a RagnarStore*

---

## Description

Inserts or updates chunks in a RagnarStore

## Usage

```
ragnar_store_insert(store, chunks)
```

```
ragnar_store_update(store, chunks)
```

## Arguments

<code>store</code>	a RagnarStore object
<code>chunks</code>	Content to insert or update. The precise input structure depends on <code>store@version</code> . See Details.

## Details

### Store Version 2

`chunks` must be `MarkdownDocumentChunks` object.

### Store Version 1

`chunks` must be a data frame containing `origin`, `hash`, and `text` columns. We first filter out chunks for which `origin` and `hash` are already in the store. If an `origin` is in the store, but with a different `hash`, we replace all of its chunks with the new chunks. Otherwise, a regular insert is performed.

This can help avoid needing to compute embeddings for chunks that are already in the store.

## Value

`store`, invisibly.

---

ragnar_store_inspect	<i>Launches the Ragnar Inspector Tool</i>
----------------------	---

---

**Description**

Launches the Ragnar Inspector Tool

**Usage**

```
ragnar_store_inspect(store, ...)
```

**Arguments**

store	A RagnarStore object that you want to inspect with the tool.
...	Passed to <code>shiny::runApp()</code> .

**Value**

NULL invisibly

---

read_as_markdown	<i>Convert files to Markdown</i>
------------------	----------------------------------

---

**Description**

Convert files to Markdown

**Usage**

```
read_as_markdown(
  path,
  ...,
  html_extract_selectors = c("main"),
  html_zap_selectors = c("nav")
)
```

**Arguments**

path	[string] A filepath or URL. Accepts a wide variety of file types, including plain text (markdown), PDF, PowerPoint, Word, Excel, images (EXIF metadata and OCR), audio (EXIF metadata and speech transcription), HTML, text-based formats (CSV, JSON, XML), ZIP files (iterates over contents), YouTube URLs, and EPUBs.
...	Passed on to <code>Markdown::convert()</code> .

**html\_extract\_selectors**

Character vector of CSS selectors. If a match for a selector is found in the document, only the matched node's contents are converted. Unmatched extract selectors have no effect.

**html\_zap\_selectors**

Character vector of CSS selectors. Elements matching these selectors will be excluded ("zapped") from the HTML document before conversion to markdown. This is useful for removing navigation bars, sidebars, headers, footers, or other unwanted elements. By default, navigation elements (nav) are excluded.

**Details****Converting HTML:**

When converting HTML, you might want to omit certain elements, like sidebars, headers, footers, etc. You can pass CSS selector strings to either extract nodes or exclude nodes during conversion.

The easiest way to make selectors is to use SelectorGadget: <https://rvest.tidyverse.org/articles/selectorgadget.html>

You can also right-click on a page and select "Inspect Element" in a browser to better understand an HTML page's structure.

For comprehensive or advanced usage of CSS selectors, consult <https://www.crummy.com/software/BeautifulSoup/bs4/doc/#css-selectors-through-the-css-property> and <https://facelessuser.github.io/soupsieve/selectors/>

**Value**

A [MarkdownDocument](#) object, which is a single string of Markdown with an @origin property.

**Examples**

```
## Not run:
# Convert HTML
md <- read_as_markdown("https://r4ds.hadley.nz/base-R.html")
md

cat_head <- \(md, n = 10) writeLines(head(strsplit(md, "\n")[[1L]], n))
cat_head(md)

## Using selector strings

# By default, this output includes the sidebar and other navigational elements
url <- "https://duckdb.org/code_of_conduct"
read_as_markdown(url) |> cat_head(15)

# To extract just the main content, use a selector
read_as_markdown(url, html_extract_selectors = "#main_content_wrap") |>
  cat_head()

# Alternative approach: zap unwanted nodes
read_as_markdown(
  url,
  html_zap_selectors = c(
```

```

      "header",      # name
      ".sidenavigation", # class
      ".searchoverlay", # class
      "#sidebar"      # ID
    )
  ) |> cat_head()

# Quarto example
read_as_markdown(
  "https://quarto.org/docs/computations/python.html",
  html_extract_selectors = "main",
  html_zap_selectors = c(
    "#quarto-sidebar",
    "#quarto-margin-sidebar",
    "header",
    "footer",
    "nav"
  )
) |> cat_head()

## Convert PDF
pdf <- file.path(R.home("doc"), "NEWS.pdf")
read_as_markdown(pdf) |> cat_head(15)
## Alternative:
# pdftools::pdf_text(pdf) |> cat_head()

# Convert images to markdown descriptions using OpenAI
jpg <- file.path(R.home("doc"), "html", "logo.jpg")
if (Sys.getenv("OPENAI_API_KEY") != "") {
  # if (xfun::is_macos()) system("brew install ffmpeg")
  reticulate::py_require("openai")
  llm_client <- reticulate::import("openai")$OpenAI()
  read_as_markdown(jpg, llm_client = llm_client, llm_model = "gpt-4.1-mini") |>
    writeLines()
  # # Description:
  # The image displays the logo of the R programming language. It features a
  # large, stylized capital letter "R" in blue, positioned prominently in the
  # center. Surrounding the "R" is a gray oval shape that is open on the right
  # side, creating a dynamic and modern appearance. The R logo is commonly
  # associated with statistical computing, data analysis, and graphical
  # representation in various scientific and professional fields.
}

# Alternative approach to image conversion:
if (
  Sys.getenv("OPENAI_API_KEY") != "" &&
  rlang::is_installed("ellmer") &&
  rlang::is_installed("magick")
) {
  chat <- ellmer::chat_openai(echo = TRUE)
  chat$chat("Describe this image", ellmer::content_image_file(jpg))
}

```



*read\_as\_markdown*

25

## End(Not run)

# Index

- \* **ragnar\_retrieve**
  - ragnar\_retrieve, 14
  - ragnar\_retrieve\_bm25, 15
  - ragnar\_retrieve\_vss, 16
- chunks\_deoverlap, 2
- chunks\_deoverlap(), 14
- commonmark::markdown\_html(), 12
- ellmer::Chat, 14, 19
- ellmer::chat\_databricks, 4
- embed\_bedrock, 3
- embed\_databricks, 4
- embed\_databricks(), 4
- embed\_google\_gemini, 4
- embed\_google\_vertex
  - (embed\_google\_gemini), 4
- embed\_lm\_studio (embed\_ollama), 6
- embed\_ollama, 6
- embed\_ollama(), 3
- embed\_openai (embed\_ollama), 6
- httr2::req\_perform\_parallel(), 13
- markdown\_chunk, 10
- markdown\_chunk(), 9
- MarkdownDocument, 8, 11, 23
- MarkdownDocument(), 9, 11
- MarkdownDocumentChunks, 9, 11
- MarkdownDocumentChunks(), 11
- ragnar\_chunks\_view, 12
- ragnar\_chunks\_view(), 11
- ragnar\_find\_links, 12
- ragnar\_read(), 12
- ragnar\_register\_tool\_retrieve, 13
- ragnar\_register\_tool\_retrieve(), 19
- ragnar\_retrieve, 14, 16, 17
- ragnar\_retrieve(), 2
- ragnar\_retrieve\_bm25, 15, 15, 17
- ragnar\_retrieve\_bm25(), 19
- ragnar\_retrieve\_vss, 15, 16, 16
- ragnar\_retrieve\_vss\_and\_bm25, 15–17
- ragnar\_store\_build\_index, 18
- ragnar\_store\_connect
  - (ragnar\_store\_create), 18
- ragnar\_store\_create, 18
- ragnar\_store\_insert, 21
- ragnar\_store\_insert(), 19
- ragnar\_store\_inspect, 22
- ragnar\_store\_update
  - (ragnar\_store\_insert), 21
- read\_as\_markdown, 22
- read\_as\_markdown(), 8
- shiny::runApp(), 22
- tibble, 2, 3
- vctrs::vec\_cast(), 19